

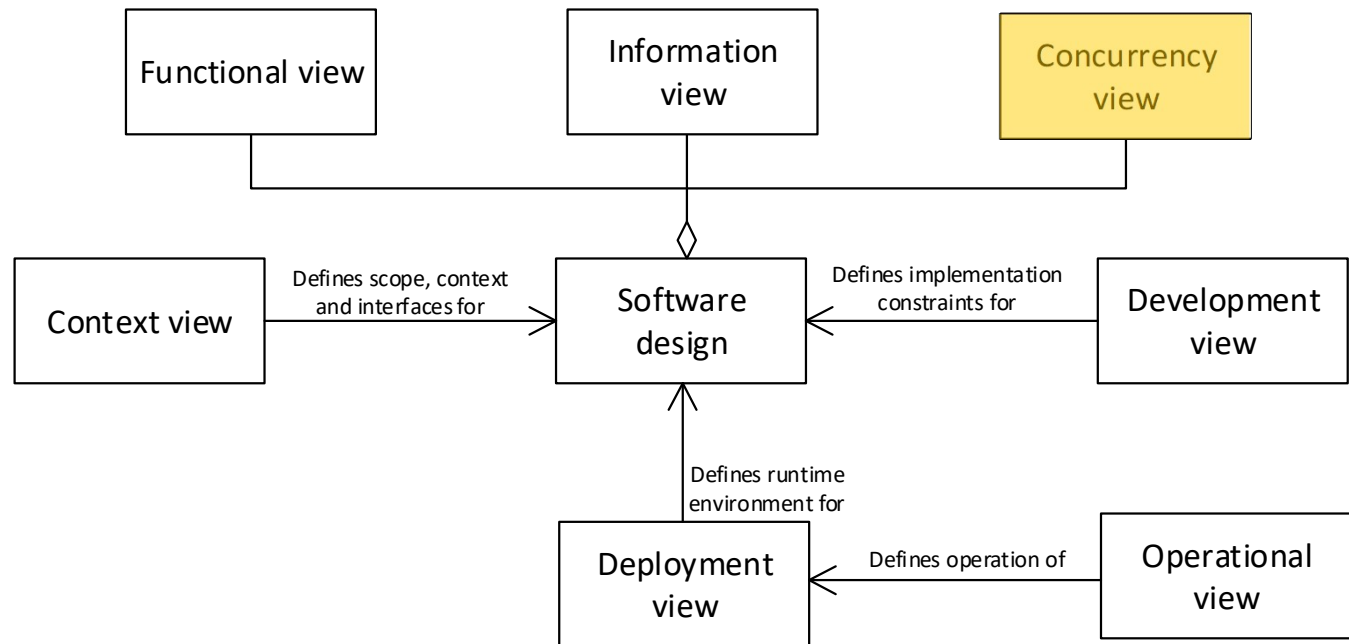


Lecture 6: concurrency viewpoint

Jan Martijn van der Werf



Viewpoint catalog

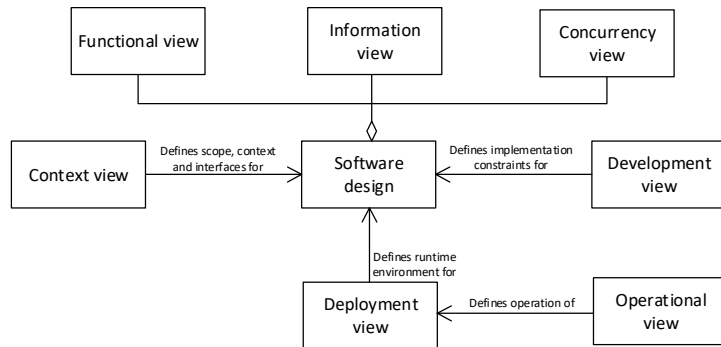


Viewpoint:

Collection of patterns, templates and conventions for constructing one type of view. It defines the stakeholders whose concerns are reflected in the viewpoint and the guidelines, principles, and template models for constructing its views



Concurrency view



- Concurrency view:
Describes the concurrency structure of the system and maps functional elements to concurrency units to clearly identify the parts of the system that can execute concurrently and how this is coordinated and controlled
- Concerns
**Task structure, mapping of functional elements to tasks,
Inter-process communication,
State management,
Synchronization and integrity,
Supporting scalability, task failure,
Startup and shutdown, re-entrancy**
- Models and views
**System level concurrency models
State models, protocol models**



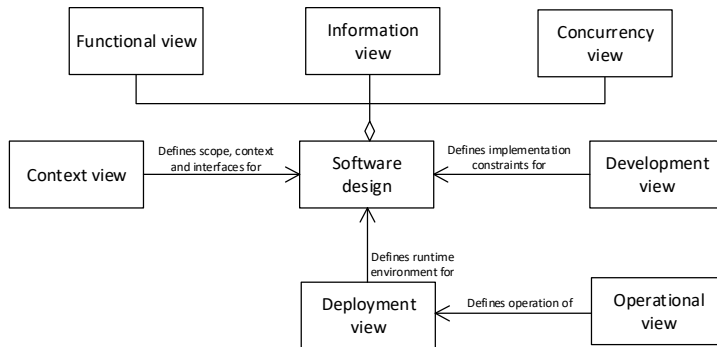
Utrecht University

About today's papers...



Concurrency view (2)

Relation with Shaw & Garlan (1995)?



- Problems and pitfalls

- Modelling the wrong concurrency**
 - Modelling the concurrency wrongly**
 - Excessive complexity**
 - Resource contention,**
 - Deadlocks, livelocks**
 - Race conditions**

- Applicability

- Systems with a number of concurrent threads of execution**
 - Component-based systems with loosely coupled elements**
 - Interorganizational systems**

We tend to illogically reason over concurrency!
It is genuinely difficult to grasp how systems are intertwined and communicate!

My position: semiformal models are insufficient to understand concurrency!






Utrecht University

Models vs. pictures

What do you see?

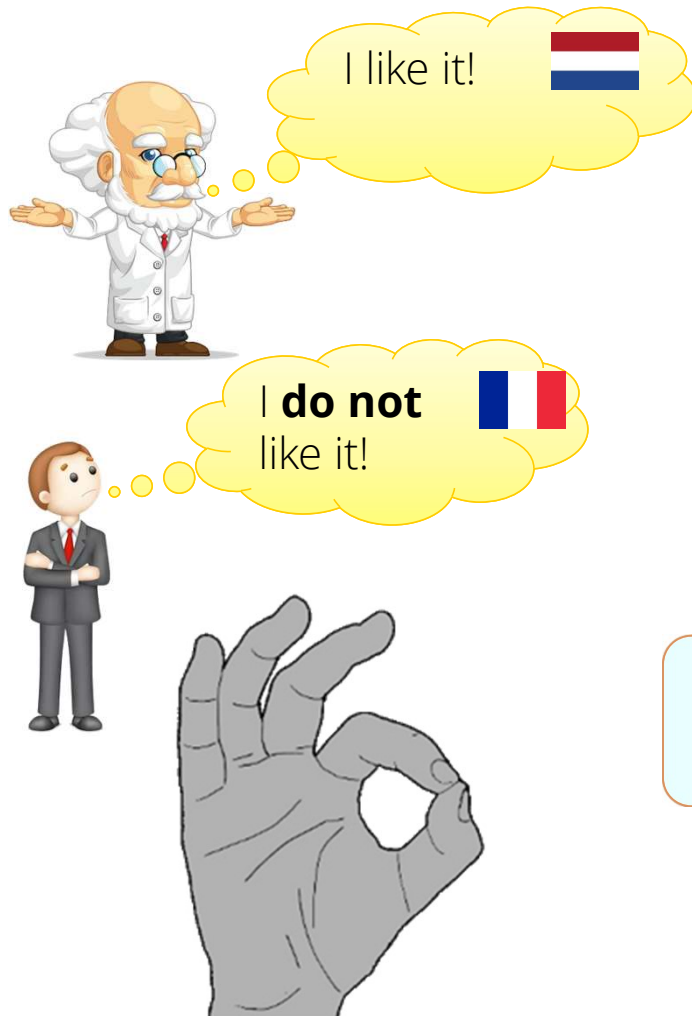


Pictures: Everybody sees something different...



Utrecht University

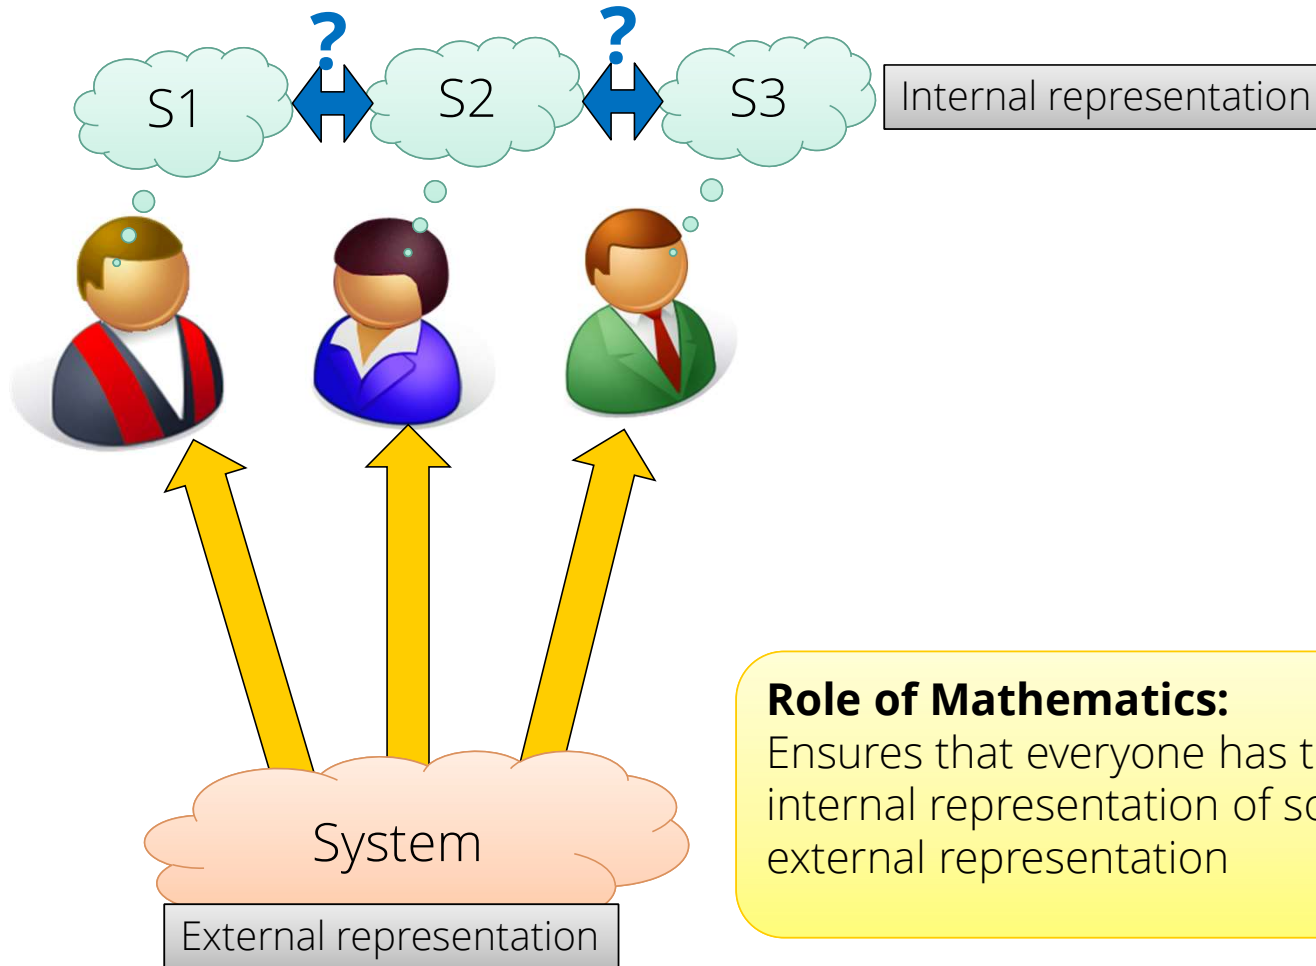
Models



- Syntax:
 - What elements are there?**
 - How are the elements related to each other?**
 - Notation:
 - How do you denote the different elements?**
 - Semantics
 - What is the (mathematical) meaning of the model?**
 - Pragmatics
 - How do you use and create the models?**
- Intention
 - What does the modeler want to express with their model?**



Intention: ensure the right internal representation!



90

ZHANG AND NORMAN





Figure 1. The theoretical framework of distributed representations. The internal representations form an internal representational space, and the external representations form an external representational space. The internal and external representational spaces together form a distributed representational space, which is a representation of the abstract task space.

Role of Mathematics:

Ensures that everyone has the same internal representation of some external representation



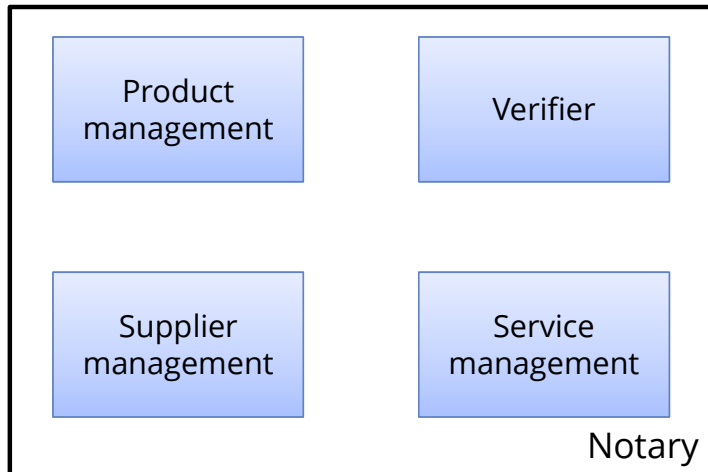


Utrecht University

Back to software architecture...

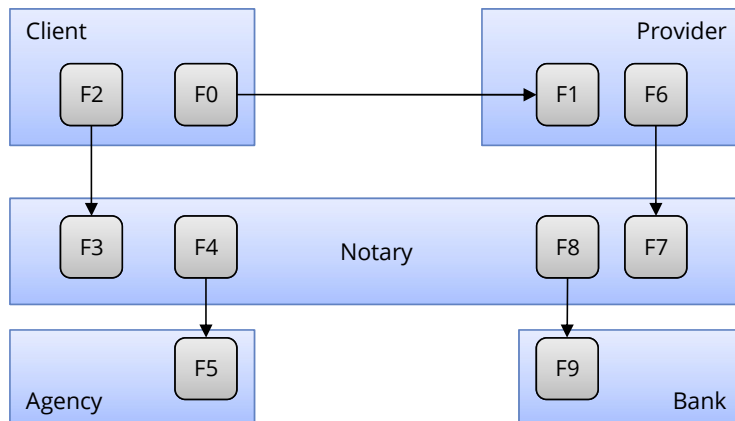


Create a functional architecture for this context:



- Notary is a platform with providers and clients.
- Providers deliver services to clients.
- Clients directly approach a provider.
- The notary is a trusted party.
- Clients need to be checked by the Notary
- Clients send their approval to a preferred provider
- Payment is only done once the client signed the bill
- Multiple payments per service possible
- Payment is done by the notary
- The notary can use several agencies to get more info upon a client's request

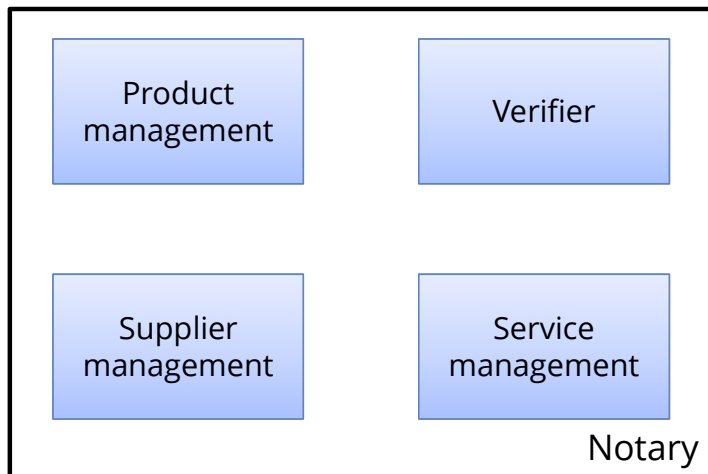
Create a functional architecture for this context:



- Notary is a platform with providers and clients.
- Providers deliver services to clients.
- Clients directly approach a provider.
- The notary is a trusted party.
- Clients need to be checked by the Notary
- Clients send their approval to a preferred provider
- Payment is only done once the client signed the bill
- Multiple payments per service possible
- Payment is done by the notary
- The notary can use several agencies to get more info upon a client's request

Create a functional architecture for this context:

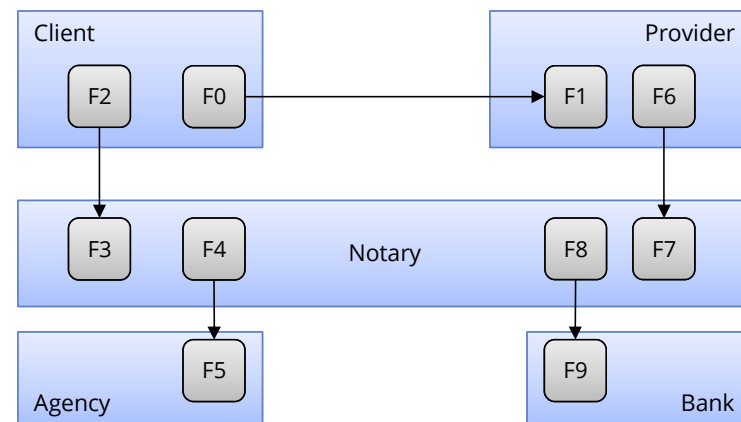
Functional architecture



From a "usage" perspective

Logical architecture

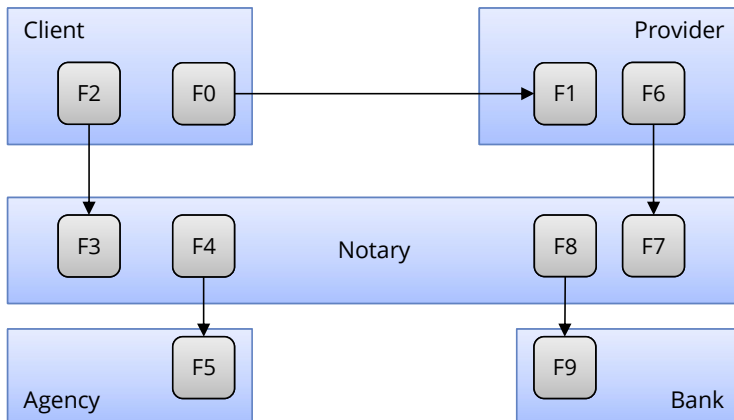
mapping



From a "resource" perspective

Here, you already see the envisioned
"services"!

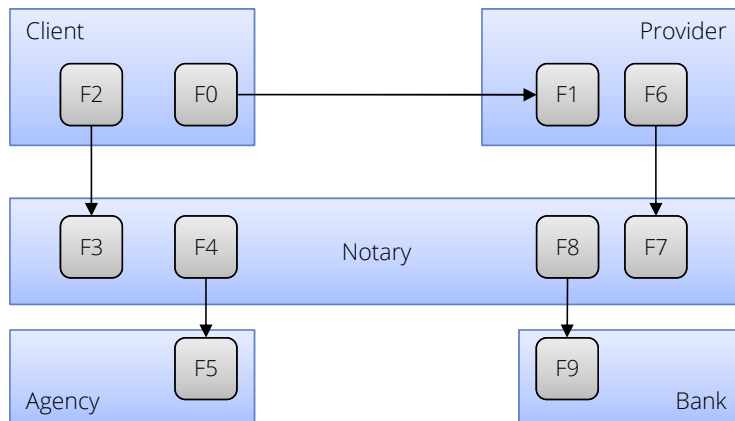
Logical architecture model



- Hierarchical structure:
Atomic and composite containers
- Atomic Container
Groups a coherent set of functions
- Composite Container
Only contains containers
- Function calls:
Always between functions
Arrow: function “converses” with another function

Remember: “just” a lines-and-boxes diagram

From logical model to concurrency...




Each system is running concurrently with all the other systems!

- Notary is a platform with providers and clients.
- Providers deliver services to clients.
- Clients directly approach a provider.
- The notary is a trusted party.
- Clients need to be checked by the Notary
- Clients send their approval to a preferred provider
- Payment is only done once the client signed the bill
- Multiple payments per service possible
- Payment is done by the notary
- The notary can use several agencies to get more info upon a client's request



Logical model and scenarios



- Scenario is a sequence of function calls

- Formal definition:

Given a logical model (C, F, h, \rightarrow) ,
a scenario is a partial order over the function calls,
i.e., $\sigma \in (\rightarrow)^*$ such that:

Functions can only start if being called before:

$$\forall 1 < i < |\sigma|: \left(\exists 1 < j < i : \pi_3(\sigma(j)) = \pi_1(\sigma(i)) \right)$$

Key:

F0: Request service	F5: Do credibility check
F1: Handle service request	F6: Request payment
F2: Request approval	F7: Check payment request
F3: Receive approval request	F8: Send payment
F4: Validate client	F9: Make payment




Given a set A , a sequence of length $n \in \mathbb{N}$ is a function $\sigma: \{1..n\} \rightarrow A$

- We write $\sigma = \langle a_1, \dots, a_n \rangle$ if $\sigma(i) = a_i$ for all $1 \leq i \leq n$.
- We denote its length by $|\sigma|$
- If $n = 0$, we call it the empty sequence, and denote it with ϵ
- The set of all finite sequences over A is denoted by A^*



Logical model and scenarios



- Scenario is a sequence of function calls

- Formal definition:

**Given a logical model (C, F, h, \rightarrow) ,
a scenario is a partial order over the function calls,
i.e., $\sigma \in (\rightarrow)^*$ such that:**

Functions can only start if being called before:

$$\forall 1 < i < |\sigma|: \left(\exists 1 < j < i : \pi_3(\sigma(j)) = \pi_1(\sigma(i)) \right)$$

Problem: scenarios can be conflicting!

Key:

F0: Request service

F1: Handle service request

F2: Request approval

F3: Receive approval request

F4: Validate client

F5: Do credibility check

F6: Request payment

F7: Check payment request

F8: Send payment

F9: Make payment




Given a set A , a sequence of length $n \in \mathbb{N}$ is a function $\sigma: \{1..n\} \rightarrow A$

- We write $\sigma = \langle a_1, \dots, a_n \rangle$ if $\sigma(i) = a_i$ for all $1 \leq i \leq n$.
- We denote its length by $|\sigma|$
- If $n = 0$, we call it the empty sequence, and denote it with ϵ
- The set of all finite sequences over A is denoted by A^*



Generalizing scenarios: Labeled transition systems



- An automaton is a 5-tuple $(Q, A, \rightarrow, s_0, \Omega)$

Q is a set of states

A is a set of observable action labels

Flow relation $\rightarrow \subseteq Q \times (A \cup \{\tau\}) \times Q$ with

τ the unobservable action, so $\tau \notin A$

$s_0 \in Q$ is the *initial state*

$\Omega \subseteq Q$ is the set of *final (accepting) states*

$Q = \{ \text{On}, \text{Off}, \text{Broken} \}$


$A = \{ a, b \},$

$\rightarrow = \{ (\text{On}, a, \text{Off}), (\text{Off}, b, \text{On}),$
 $(\text{Off}, \tau, \text{Broken}), (\text{On}, \tau, \text{Broken}) \}$

$s_0 = \text{On}$

$\Omega = \{ \text{Off}, \text{Broken} \}$

Modeling Label Transition Systems in LTSA



- States in capital
- Actions in small letters
- List all transitions per state:

Note: silent action cannot be implemented directly!
Choose a name not in A.

```

LAMP = ON,
ON   = ( a -> OFF
        | t -> BROKEN
        ),
OFF  = ( b -> ON
        | t -> BROKEN
        ).

|| SYS = (LAMP) \ {t}.
    
```

Transition symbol

Comma between states


Dot at the end

Non deterministic choice

SYS is a composition

On LAMP, hide action t

Define an LTS for F0 in LTSA





Key:

F0: Request service	F5: Do credibility check
F1: Handle service request	F6: Request payment
F2: Request approval	F7: Check payment request
F3: Receive approval request	F8: Send payment
F4: Validate client	F9: Make payment

- The client asks for permission. Upon receiving the permission, (s)he asks for a service of the provider, and either accepts or denies the offer received by the provider. Upon accepting, the provider offers the service, and regularly sends an invoice, which needs to be signed and returned by the client. Once denied, either (s)he stops, or (s)he asks for a better offer.
- In case the client does not receive permission, (s)he tries again.



Interface automata (Also referred to as Finite State Processes)



Question mark:
receives a message
from the environment



Exclamation mark:
sends a message to
the environment

- Actions of LTS divided in three sets:
 - Internal action**
 - Sending action: denoted with an exclamation mark**
 - Receiving action: denoted with a question mark**
- Synchronization:
 - Sync is a function: if $a!$ and $a?$ are both enabled in the two LTSs, they synchronize in action a**



Composing interface automata

Composition of interface automata





Composing interface automata

Composition of interface automata

Let $L_1 = (S_1, A_1, \rightarrow_1, s_0^1, \Omega_1)$ **and** $L_2 = (S_2, A_2, \rightarrow_2, s_0^2, \Omega_2)$,

Let $R \subseteq A$ **be a set of actions to synchronize on**

i.e. $\forall a \in R: (a? \in S_1 \wedge a? \in S_2) \vee (a! \in S_1 \wedge a! \in S_2)$

Define $R^+ = \bigcup_{a \in R} \{a!, a?\}$

Then: $L_1 \otimes_R L_2 = (S, A, \rightarrow, s_0, \Omega)$ **with**

$$S = S_1 \times S_2$$

$$A = (R \cup A_1 \cup A_2) \setminus R^+$$

$$\rightarrow = \{ ((s_1, s_2), a, (s'_1, s'_2)) \mid (s_1, a, s'_1) \in \rightarrow_1 \wedge a \in A_1 \setminus R^+ \}$$

$$\cup \{ ((s_1, s_2), a, (s'_1, s'_2)) \mid (s_2, a, s'_2) \in \rightarrow_2 \wedge a \in A_2 \setminus R^+ \}$$



$$\cup \{ ((s_1, s_2), a, (s'_1, s'_2)) \mid (s_1, a?, s'_1) \in \rightarrow_1, (s_2, a!, s'_2) \in \rightarrow_2, a \in R \}$$

$$\cup \{ ((s_1, s_2), a, (s'_1, s'_2)) \mid (s_1, a!, s'_1) \in \rightarrow_1, (s_2, a?, s'_2) \in \rightarrow_2, a \in R \}$$

$$s_0 = (s_0^1, s_0^2)$$

$$\Omega = \Omega_1 \times \Omega_2$$




Notice: this is interleaving semantics!



ReqEng: In FSP: $(a \rightarrow S) \parallel (b \rightarrow T) \equiv (a \rightarrow (S' \parallel (b \rightarrow T))) \mid b \rightarrow ((a \rightarrow S) \parallel T))$



Composing interface automata



$$R = \{a, b\}$$





Composing interface automata in LTSA

- Define each system separately

Prefix a send action with an s



Prefix a receive action with an r



```
LAMP = ON,  
ON   = ( rOff -> OFF | t -> BROKEN),  
OFF  = ( rOn  -> ON  | t -> BROKEN).  
  
SWITCH = UP,  
UP     = ( sDown -> DOWN),  
DOWN  = ( sUp   -> UP).
```




Composing interface automata in LTSA



- Define each system separately
Prefix a send action with an s
Prefix a receive action with an r
- Compose the systems (||)

```
LAMP = ON,
ON    = ( rOff -> OFF | t -> BROKEN ),
OFF   = ( rOn  -> ON  | t -> BROKEN ).


SWITCH = UP,
UP      = ( sDown -> DOWN ),
DOWN    = ( sUp   -> UP ).

||SYS = ( LAMP || SWITCH ).
```

SYS is a composition of lamp with switch



Composing interface automata in LTSA



- Define each system separately
Prefix a send action with an s
Prefix a receive action with an r
- Compose the systems ($||$)
- Specify synchronization relation with relabelling:
rename the actions to create the synchronization:
 $/\{s\text{SendAction}/r\text{ReceiveAction}\}$

```
LAMP = ON,
ON    = ( rOff -> OFF | t -> BROKEN),
OFF   = ( rOn  -> ON  | t -> BROKEN).




SWITCH = UP,
UP     = ( sDown -> DOWN),
DOWN  = ( sUp   -> UP).

||SYS = ( LAMP || SWITCH )/{sDown/rOn, sUp/rOff}.
```

/ denotes a renaming action!
s/r renames an r into s



Composing interface automata






$$R = \{a, b\}$$

Notice: once the lamp is broken, it is not possible to use the switch anymore! This is due to the fact that the actions of the switch need to be synchronized with the lamp. As it does not react on any messages anymore, the system is in **deadlock**.

Deadlock: a state with no outgoing actions. Possible repair: add a "self-loop" on state "bro" labeled *a?* and *b?*



Composing interface automata




$$R = \{a, b\}$$

Notice: once the lamp is broken, it is not possible to use the switch anymore! This is due to the fact that the actions of the switch need to be synchronized with the lamp. As it does not react on any messages anymore, the system is in **deadlock**.

Deadlock: a state with no outgoing actions. Possible repair: add a "self-loop" on state "bro" labeled *a?* and *b?*

Define the following LTS for F1 in LTSA, and compose it with F0



- Upon receiving a request for an offer from a client, the provider creates and sends an offer. If the offer is rejected, the provider will create a new offer, and send it.
- Once the offer is accepted, the service is provided, and on a regular basis an invoice is sent, which is signed and returned by the client. This signed bill is then passed to the administration.

Key:


F0: Request service	F5: Do credibility check
F1: Handle service request	F6: Request payment
F2: Request approval	F7: Check payment request
F3: Receive approval request	F8: Send payment
F4: Validate client	F9: Make payment

What are your observations?

Designing the individual systems is “easy”. Their composition is exponentially more difficult!



Communication mechanisms



- Synchronous communication
You have to know in which state the other is!
- Asynchronous communication
Send messages to the other, the state of the other is unknown!





FSP, Interface automata only allow for synchronous communication!

Key:

F0: Request service	F5: Do credibility check
F1: Handle service request	F6: Request payment
F2: Request approval	F7: Check payment request
F3: Receive approval request	F8: Send payment
F4: Validate client	F9: Make payment





Simulating asynchronous behavior with buffers





I/O automata



- Message queue to delay messages
Queue per automaton / global queue
Messages are handled in order.
Capacity of the queue?




Utrecht University

Petri nets



Labeled Transition Systems have global states!



$R = \{a, b\}$

State resembles:


- State of switch
- State of light



How can we make advantage of this split state?

- Instead of global states: distributed states!
- Transitions: state changes are local
check locally if a transition can fire:
transition firing changes the state locally

Labeled Transition Systems have global states!



$R = \{a, b\}$

State resembles:

- State of switch
- State of light



How can we make advantage of this split state?

- Instead of global states: distributed states!



Places: indicate a condition of the system


Tokens: represents elementary entities of interest

State: the configuration of tokens over the different places

- Transitions: state changes are local
check locally if a transition can fire:
transition firing changes the state locally



Petri nets: an informal introduction



- Instead of global states: distributed states!



Places: indicate a condition of the system


Tokens: represents elementary entities of interest

State: the configuration of tokens over the different places

- Transitions: state changes are local
check locally if a transition can fire:
transition firing changes the state locally



Petri nets: an informal introduction



- Instead of global states: distributed states!



Places: indicate a condition of the system

Tokens: represents elementary entities of interest

State: the configuration of tokens over the different places

- Transitions: state changes are local



check locally if a transition can fire:


transition firing changes the state locally



Arcs are always from a place to a transition, or
From a transition to a place



Petri nets: an informal introduction



- Instead of global states: distributed states!



Places: indicate a condition of the system

Tokens: represents elementary entities of interest

Marking: configuration of tokens over the different places

- Transitions: state changes are local




Input places: check locally if a transition can fire:

Output places: transition firing changes the state locally



- Transition enabled:
all input places have at least 1 token



Petri nets: an informal introduction




output places of transition t

- Instead of global states: distributed states!
-  **Places:** indicate a condition of the system
Tokens: represent elementary entities of interest
Marking: configuration of tokens over the different places
- Transitions: state changes are local
 **Input places:** check locally if a transition can fire:
Output places: transition firing changes the state locally
- Transition enabled:
all input places have at least 1 token
- Transition firing (executing the transition):
Consumes 1 token from each input place
Produces 1 token in each output place



Petri nets: an informal introduction



There is no law of token preservation!

- Instead of global states: distributed states!



Places: indicate a condition of the system

Tokens: represent elementary entities of interest

Marking: configuration of tokens over the different places

- Transitions: state changes are local




Input places: check locally if a transition can fire:


Output places: transition firing changes the state locally

- Transition enabled:
all input places have at least 1 token
- Transition firing (executing the transition):
Consume 1 token from each input place
Produce 1 token in each output place

Petri nets: an informal introduction



- Instead of global states: distributed states!



Places: indicate a condition of the system

Tokens: represent elementary entities of interest

Marking: configuration of tokens over the different places

- Transitions: state changes are local



Input places: check locally if a transition can fire:



Output places: transition firing changes the state locally

- Transition enabled:
all input places have at least 1 token
- Transition firing (executing the transition):
Consume 1 token from each input place
Produce 1 token in each output place

There is no law of token preservation!




Petri nets: behavior in terms of LTSs



these two models describe the same behavior!



Petri nets: formal definition



$N = ((P, T, F), m_0)$ with:

$$P = \{p, q, r, s\}$$

$$T = \{t, u, v\}$$

$$F = \{(p, t), (r, t), (t, q), (q, u), \\ (u, s), (s, v), (v, r), (v, p)\}$$

$$m_0 = [p, r^2]$$


(reads as: $[p^1, r^2, q^0, s^0]$)

- A Petri net is a 4-tuple $N = ((P, T, F), m_0)$ with:
 - P : a (finite) set of places
 - T : a (finite) set of transitions
 - P and T are disjoint ($P \cap T = \emptyset$)
 - F is the flow relation that defines the arcs
- $$F \subseteq (P \times T) \cup (T \times P)$$
- m_0 is the initial marking, gives the tokens per place
- $$m_0: P \rightarrow \mathbb{N}$$




Petri nets vs. LTSs

All input places always have a token
(empty set, so always true!)




Always enabled!



Infinite sequence



Infinite Automata!

Back to our lamps: now as Petri nets







Back to our lamps: now as Petri nets



Lamp





Switch








Back to our lamps: now as Petri nets




Lamp



Switch



Lamp



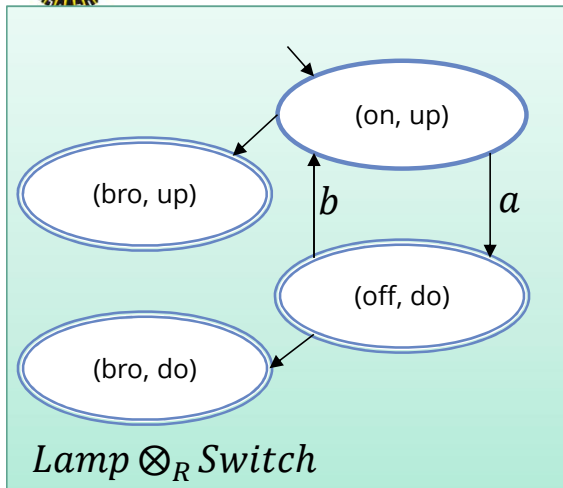


Switch




Back to our lamps: now as Petri nets

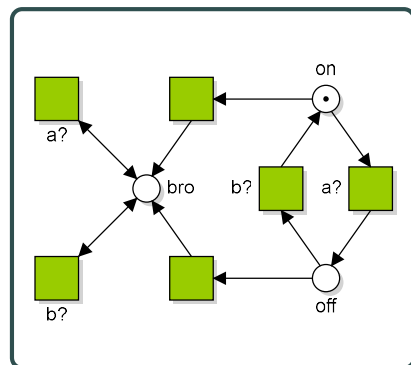

Synchronous communication




Lamp



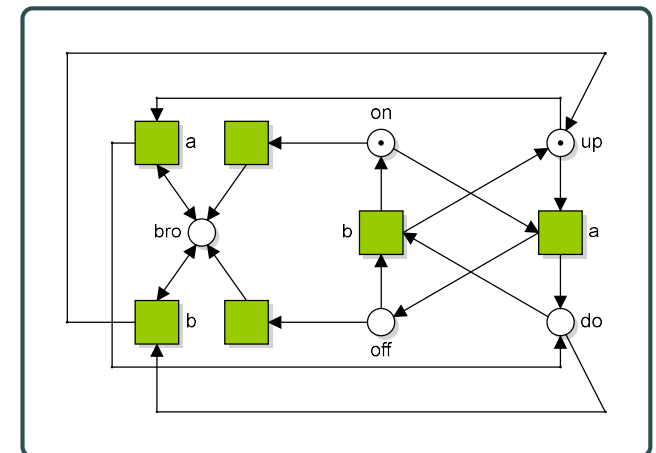
Switch



Lamp




Switch






Back to our lamps: now as Petri nets Asynchronous communication!






Transition receives a message by consuming a token from "message place" b



Transition sends a message by placing a token in the "message place" b




Back to our lamps: now as Petri nets Asynchronous communication!





Petri nets and communication mechanisms



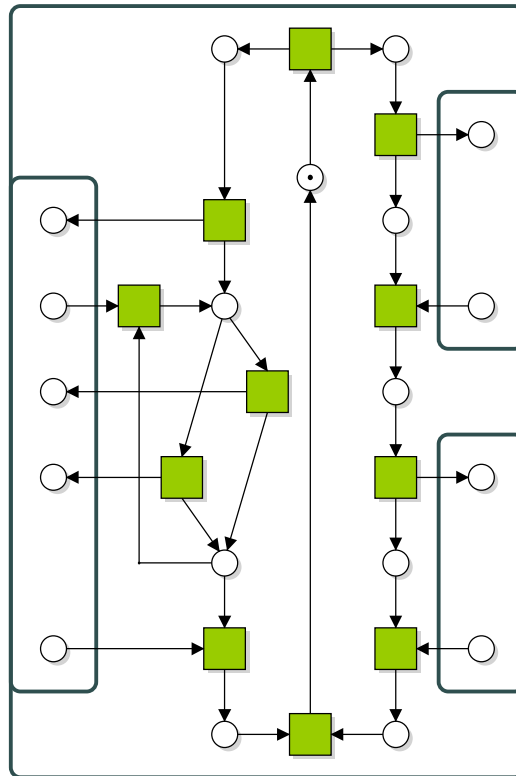

- Synchronous communication
Transitions that consume from multiple places
- Asynchronous communication
Places resemble pools of messages to be handled
Random access of messages in these pools



Both communication mechanisms in a single, **graphical** formalism!



Teaser for next week:



- Modelling communication between components
- Different perspectives on message handling



Utrecht University



For the remainder of today:

- Work on the assignment
- Thursday: first presentation session!
 - ➔ **No need for a complete presentation**
 - ➔ **Schedule will be made available Thursday morning**



Utrecht University

DISCLAIMER

The information in this presentation has been compiled with the utmost care,
but no rights can be derived from its contents.

© Utrecht University