



# *Cloud Architecture*

**Nishant Saurabh**



## Agenda for today

- 13:15 – 14:15: Cloud Architecture
- 14:30 – 14:45: Lab Assignment: I
- 14:45 – 15:45: Lab-I Assignment time
- 16:00 – 16:45: Assignment time
- 16:45 – 17:00: Wrap-up





Utrecht University

# *Service-oriented Architecture?*



## *Definition 1*

*Service –oriented architectures (SOAs) are way of developing distributed systems where system components are stand-alone services , executing on geographically distributed servers*




## Definition 2

*A paradigm for organizing and utilizing distributed capabilities that may be under the control of different ownership domains. It provides uniform means to offer, discover, interact with and use capabilities to produce desired effects consistent with measurable preconditions and expectations*

**OASIS (Organization for the Advancement of Structured Information Standards)**



# Components of SOA




- Services
  - Loosely-coupled reusable software component**
  - Encapsulates discrete functionalities**
  - Distributed and programmatically accessed**
- Messages
  - Exchange of information**
- Meta-data
  - Service description**
  - Service interface**
  - Service metadata**



## SOA characteristics


- Integration of SOA architecture
- Application
- Scope
- Dependency
- Stateful vs stateless
- Loose coupling
- Reusability
- Platform and implementation independent
- Discoverable





## SOA considerations?

- Computational penalties
- Extra layers
- Higher communication latency
- Supporting native applications
- Service granularity (reusability or performance)
- Fault tolerance (partial or full failure)
- Service agreements (performance and cost)
- Governance (how to manage and orchestrate)







Utrecht University

# *Cloud Architecture*



# You have house to rent



- What does the tenant want?
- What can you offer ?



# You have house to rent




- What does the tenant want?
- What can you offer ?
- Is it affordable ?
- Is it spacious?
- Will I be disturbed by outsiders?
- Will energy cost be billed separately?



# You have a computer to rent

- What does the tenant want?
- What can you offer ?
- Is it affordable to rent?
- Is there enough CPU/Memory/Disk?
- Is network connection efficient?
- Do I pay for what I use?



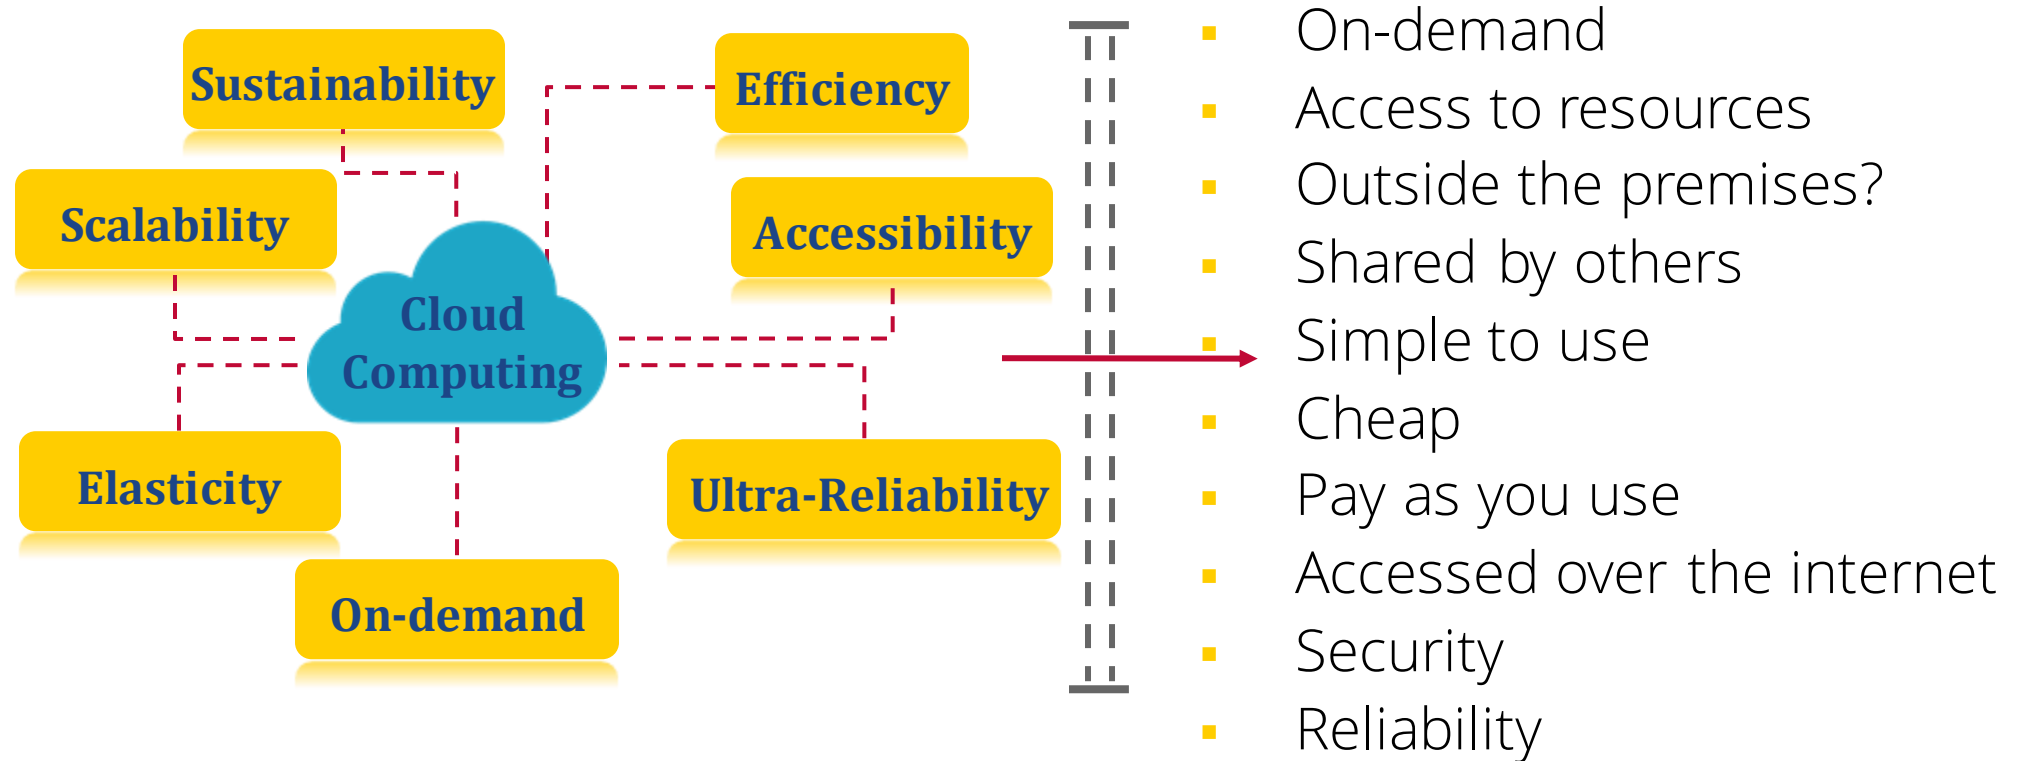


## *NIST Definition of Cloud*

*A model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g. networks, servers, storage, applications and services) that can rapidly provisioned and released with minimal management effort or service provider interaction*




# What makes Cloud Ubiquitous?





# What makes Cloud Ubiquitous?



- On-demand self service  
**Unilaterally provision computing capabilities without requiring human interactions**
- Broad network access  
**Available over the network**  
**Heterogenous thin or thick client platforms (mobile phones, laptops)**
- Resource pooling  
**Compute resources serve multiple users**  
**Multi-tenant model**
- Metering capability  
**Monitoring services**  
**Automatic control and optimization**
- Rapid elasticity  
**Scale out**  
**Scale in**

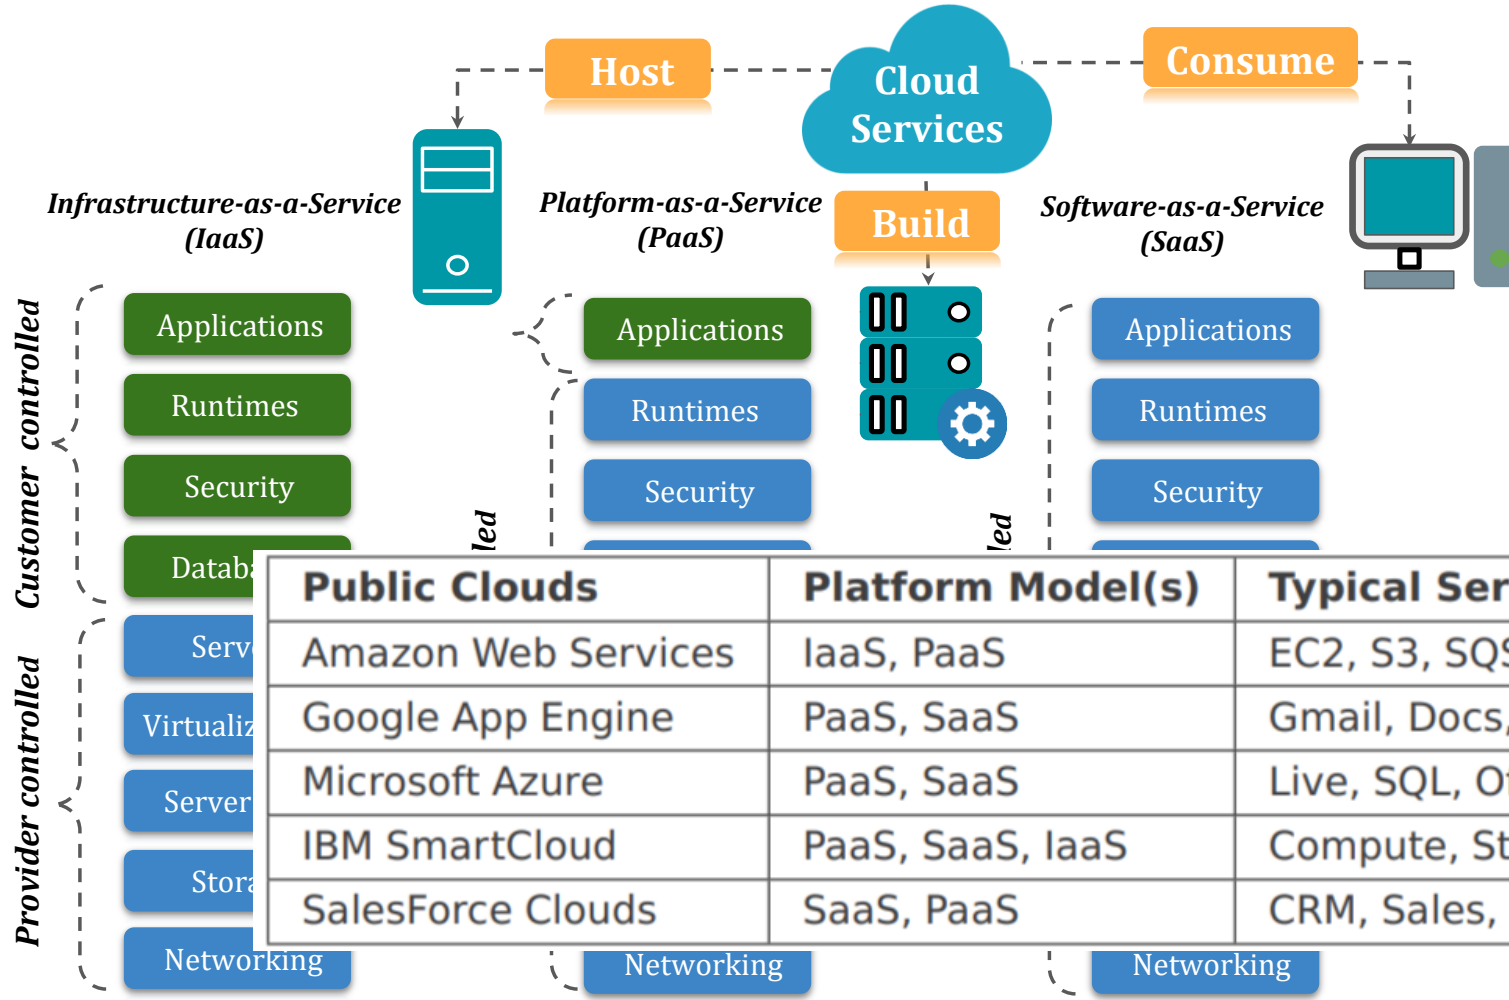


*What are Cloud service models ?*





# Cloud service models



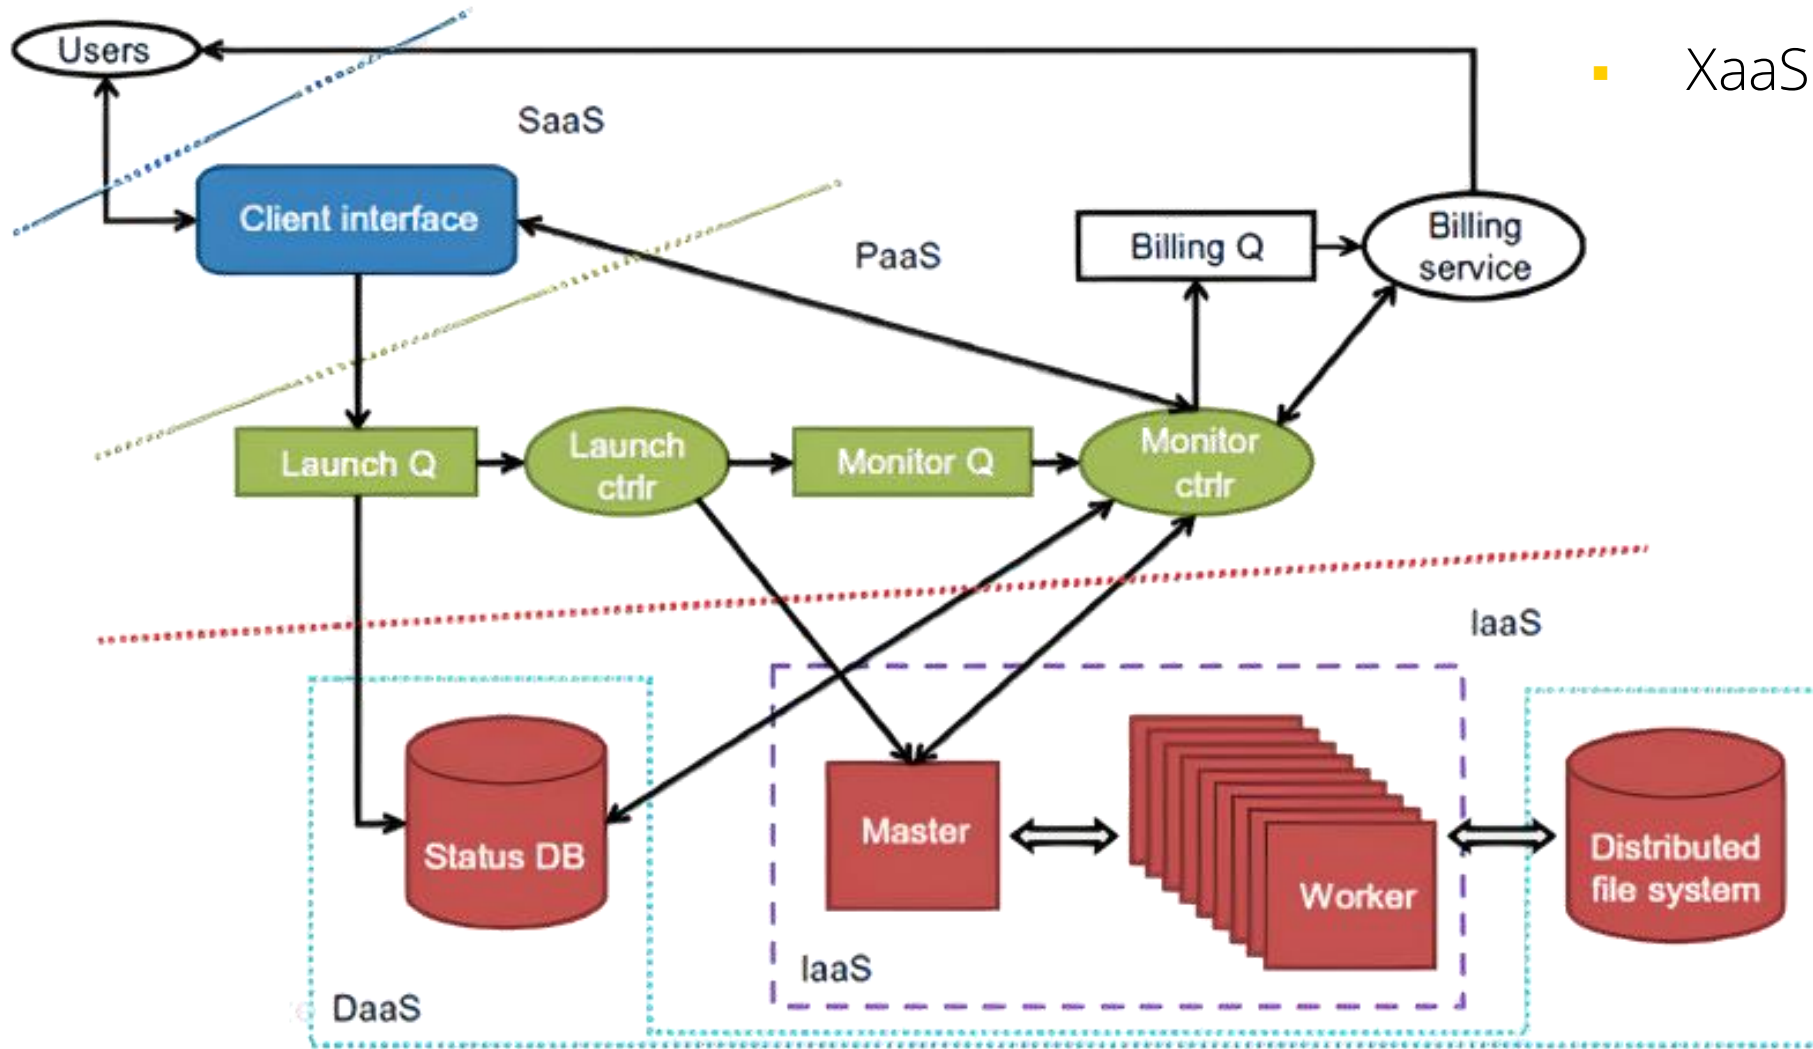
- Software-as-a-Service
- Platform-as-a-Service
- Infrastructure-as-a-Service

Loss of Control



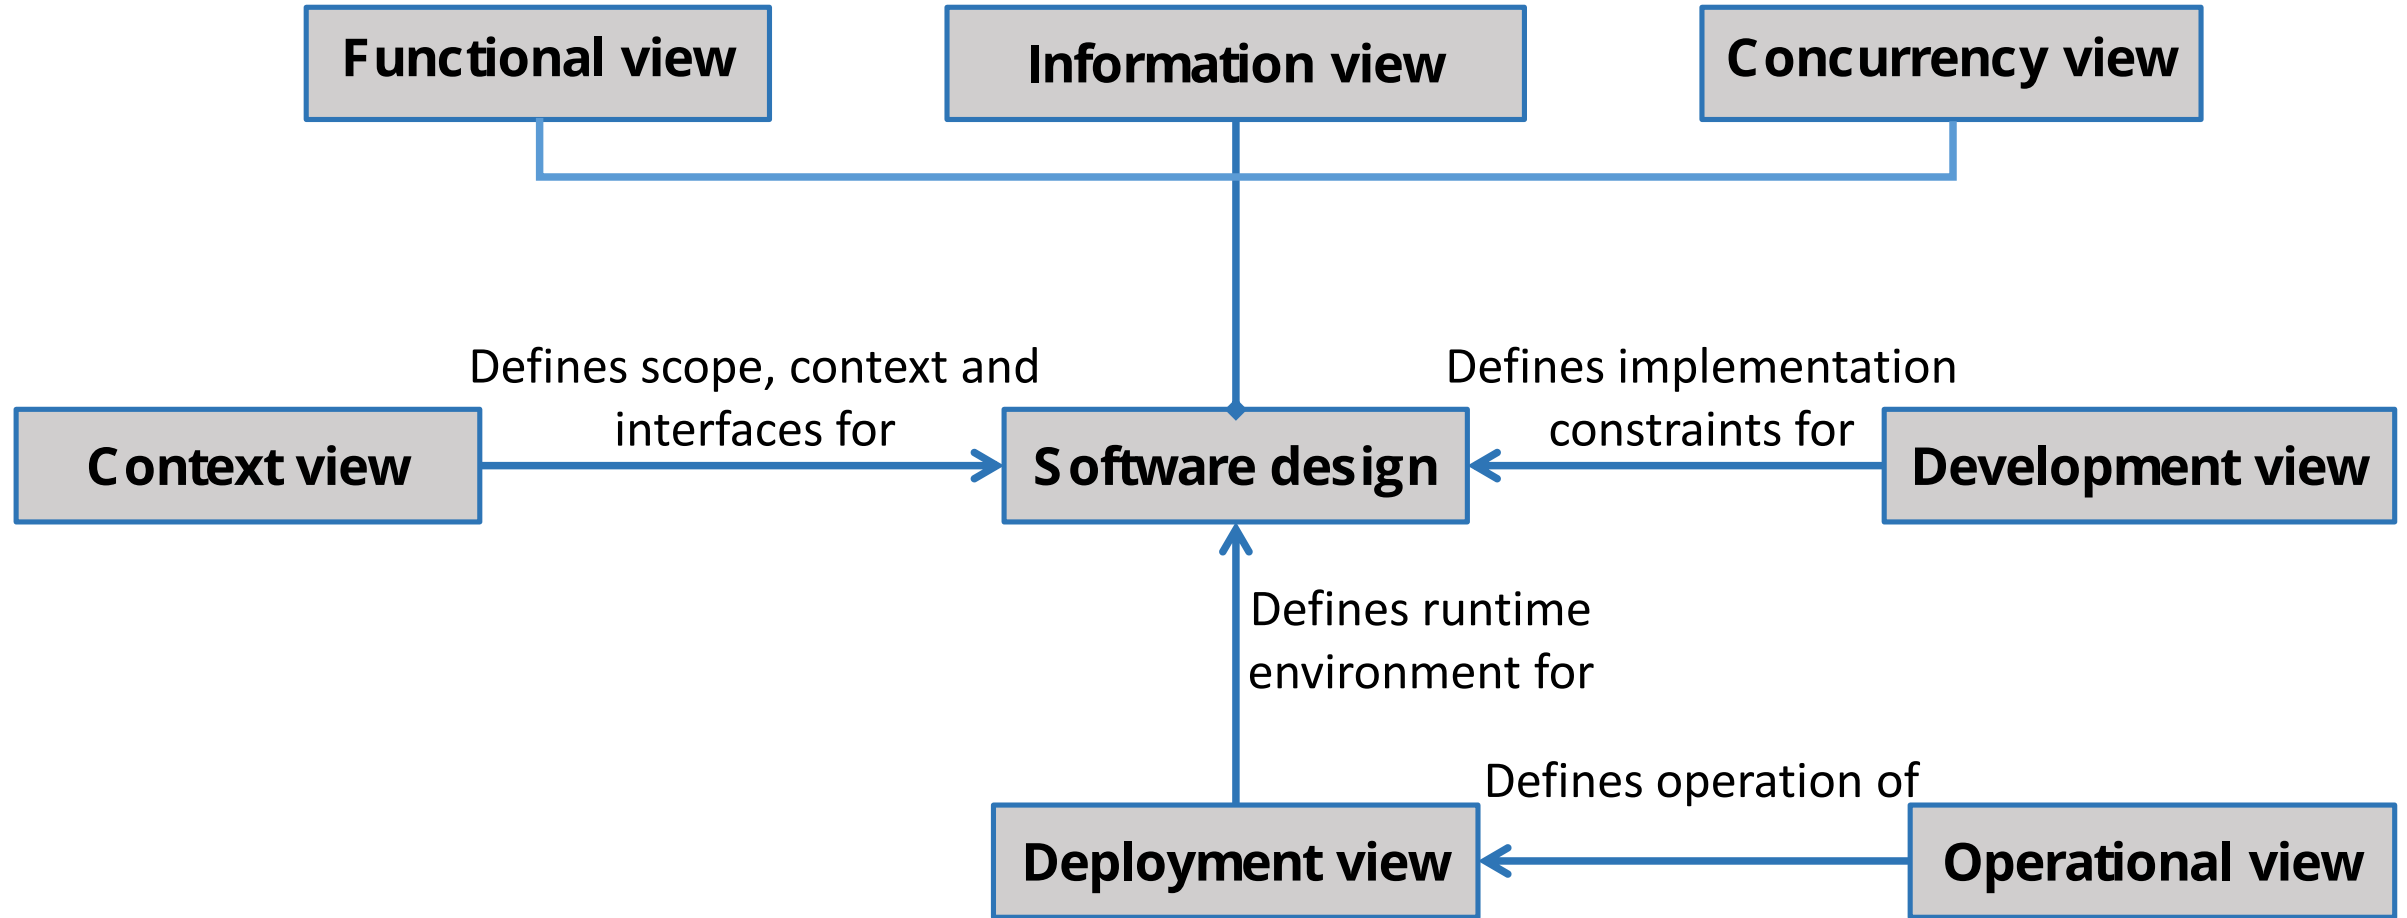
# XaaS: Everything-as-Service

- XaaS architectural hierarchy





# Where does Cloud fit?





# Cloud deployment models

- Public Cloud  
**Available to all**  
**Example: AWS, GCP**
- Private Cloud  
**On-premise**  
**Example: OpenStack**
- Hybrid Cloud  
**Private + Public**  
**E.g. Rackspace Cloud**
- Community Cloud  
**Shared by organisations with similar goals**  
**Example: salesforce**



Cloud Architecture: from chaotic to complex to complicated... ?



*Why do we need Cloud?*



# Economies of Scale



- Large data centers cheaper to operate
- Large means 100,000+ servers
- Small means less than 10,000 servers



# Economies of Scale



- Large data centers cheaper to operate
- Large means 100,000+ servers
- Small means less than 10,000 servers

Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?





# Economies of Scale



- Cost of power  
**Lower per server**
- Infrastructure labor costs  
**Administrator services more than 1000 servers**
- Security and reliability investment  
**Disaster recovery**
- Hardware costs  
**Large scale hardware purchase**

Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?





# Utilization



- Collocated services and applications
- Variations in workload
- Time of the day, year
  - Resource usage patterns**
  - Uncertainty (e.g. news events, marketing events)**

Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?



# Multitenancy

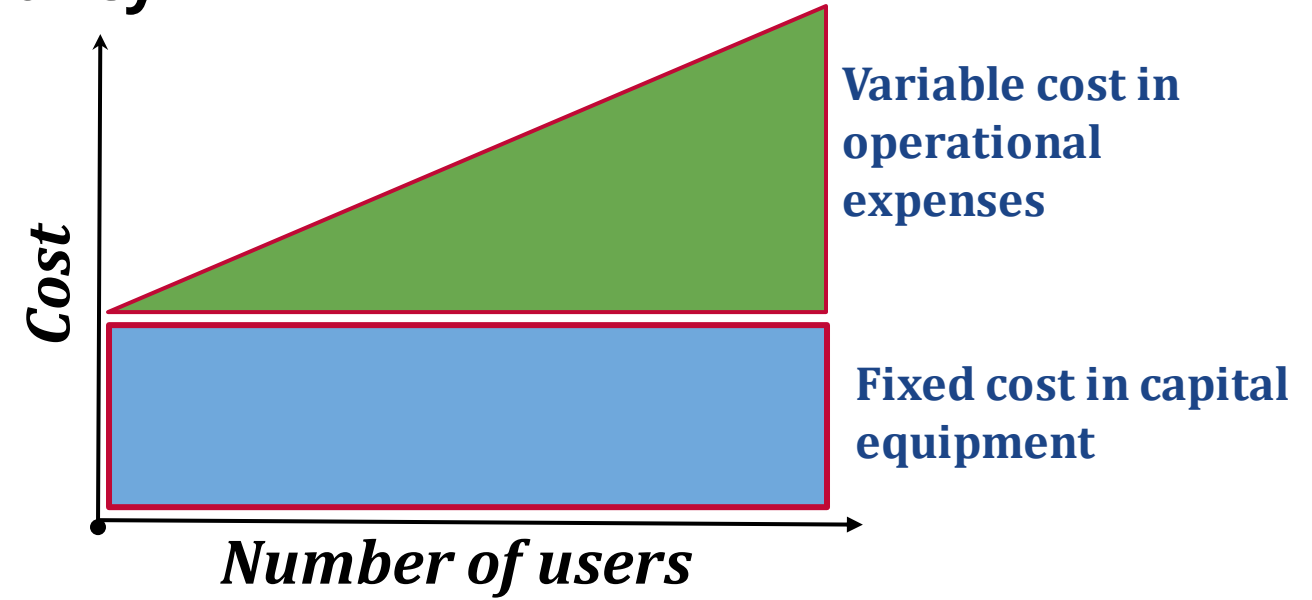


- Salesforce  
**Single application for multiple users**
- Reduced costs  
**help desk support**  
**upgradation costs**  
**software development and maintenance costs**

Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?




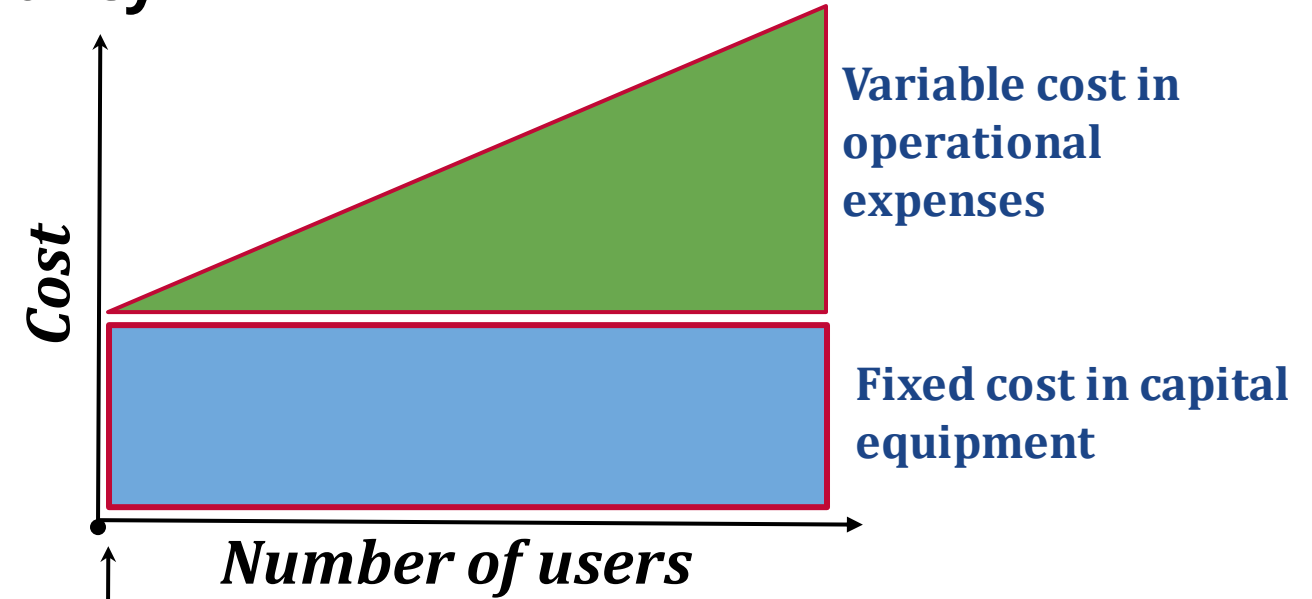
# Multitenancy



Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?



# Multitenancy



Microsoft datacenter: 11.5 times the size of a football field ? Still economic ?



# *Virtualization: VMs to Serverless*



# Virtualization

- Create a software-based or virtual representation of applications, servers, storage and networks to reduce IT expenses while boosting efficiency and agility.





# Virtualization

- Create a software-based or virtual representation of applications, servers, storage and networks to reduce IT expenses while boosting efficiency and agility.

Virtualization != Cloud computing?



# Virtualization abstraction levels

| Level of Virtualization              | Functional Description                                  | Example Packages            | Merits, App Flexibility/ Isolation, Implementation Complexity                       |
|--------------------------------------|---|-----------------------------|---|
| <b>Instruction Set Architecture</b>  | Emulation of a guest ISA by host                        | Dynamo, Bird, Bochs, Crusoe | Low performance, high app flexibility, median complexity and isolation              |
| <b>Hardware-Level Virtualization</b> | Virtualization on top of bare-metal hardware            | XEN, VMWare, Virtual PC     | High performance and complexity, median app flexibility, and good app isolation     |
| <b>Operating System Level</b>        | Isolated containers of user app with isolated resources | Docker Engine, Jail, FVM    | Highest performance, low app flexibility and best isolation, and average complexity |
| <b>Run-Time Library Level</b>        | Creating VM via run-time library through API hooks      | Wine, vCUDA, WABI, LxRun    | Average performance, low app flexibility and isolation, and low complexity          |
| <b>User Application Level</b>        | Deploy HLL VMs at user app level                        | JVM, .NET CLR, Panot        | Low performance and app flexibility, very high complexity and app isolation         |



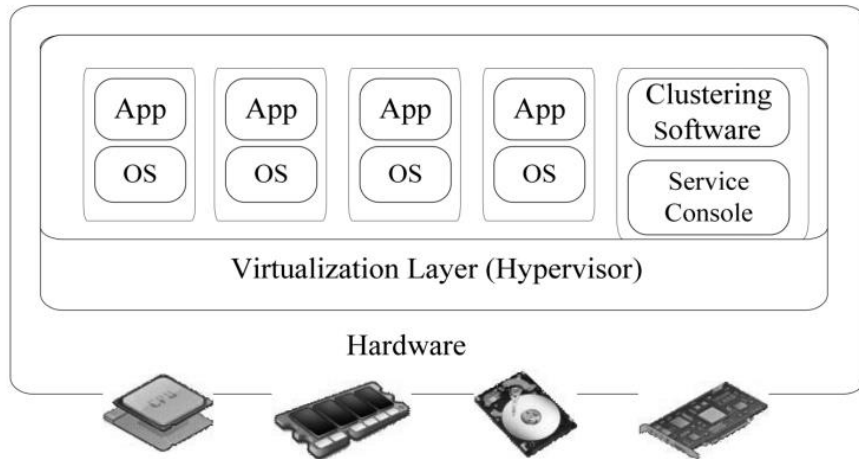


# *Hardware and OS-level virtualization*



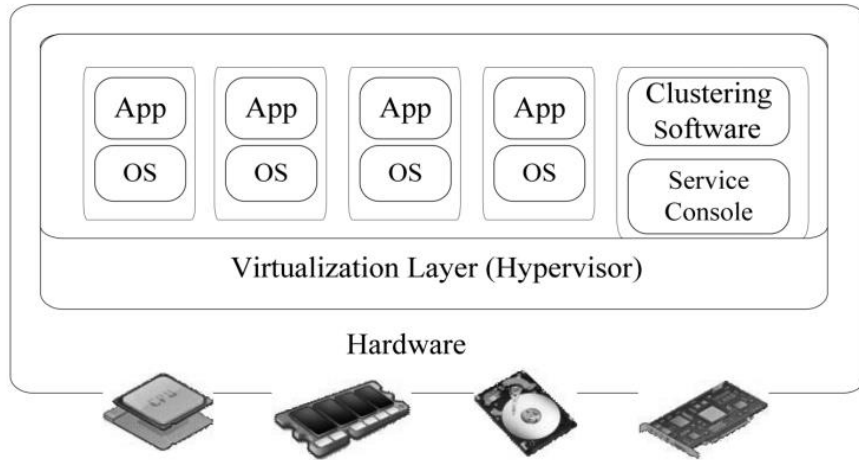
# Hardware-level virtualization

- Virtual hardware environment for VMs  
**Xen, KVM, Hyper-V, Virtual Box**
- Virtual Machine (VM)  
**Tightly isolated software container with an OS and application inside, stateful in nature**
- Hypervisor  
**Thin layer of software decouples VMs from host and dynamically allocates computing resources to each VM**





# Hardware-level virtualization



- Virtual hardware environment for VMs  
**Xen, KVM, Hyper-V, Virtual Box**
- Virtual Machine (VM)  
**Tightly isolated software container with an OS and application inside, stateful in nature**
- Hypervisor  
**Thin layer of software decouples VMs from host and dynamically allocates computing resources to each VM**

Drawbacks: Long boot delay, Takes upto tens of GBs, VM sprawl issues


Saurabh, Nishant, et al. "Semantics-aware virtual machine image management in IaaS clouds." 2019 IEEE International Parallel and Distributed Processing Symposium (IPDPS). IEEE, 2019.


Saurabh, Nishant, et al. "Expelliarmus: Semantic-centric virtual machine image management in IaaS Clouds." *Journal of Parallel and Distributed Computing* 146 (2020): 107-121.



# OS-level virtualization

- Containers  
**Creates isolated containers on a single physical server or a VM to utilize hardware and software in datacenters**  
**E.g. Dockers**





# OS-level virtualization

- Containers
  - Creates isolated containers on a single physical server or a VM to utilize hardware and software in datacenters**
  - E.g. Dockers, stateless in nature**
- Pods
  - Group of related containers**
  - Used for deploying, managing and scaling containers**


Advantages: Lower startup time compared to VMs  
Drawbacks: Performance interference, Container image sizes are growing upto few GBs, more start up time



# Serverless

## Event-driven

- Datastores (e.g. S3)
- Message Queues
- Applications
- Sensors
- Scheduled tasks (via Alarms)



- Function as a Service (FaaS)
  - Light weight functions**
  - Can be stored in a repository**
  - Runtime system to run any kind of code**
  - Java, Node.js, Python etc.**
- Place functions in containers
- Real pay as you go model
  - Use resources only when you run functions**
- E.g. AWS Lambda, Azure functions



# VMs/Containers vs. Serverless

- VM and containers reserve resources without a load  
**Can be costly**
- Serverless uses resources only when a function is invoked  
**Provider waits for requests**
- Serverless computing  
**Also serverless DB, etc.**
- Disadvantages  
**Cold start problem**



Utrecht University

# *Architecting in a Cloud environment*





# Architecting in a Cloud Environment



- Quality attributes that are different in a cloud
  - Security and Consistency**
  - Performance**
  - Availability**



# Security and Consistency

- Multi-tenancy concerns over non-cloud environments
  - Inadvertent information sharing : shared use of resources**
  - Virtual machine escape : breaking the hypervisor**
  - Side channel attacks: monitoring cache**
  - Denial of service attacks**
- Some systems require:
  - state management, time coordination and data coordination**
- You need to consider risks when deciding what applications to host in the cloud
- Also which services need consistency measures



# Performance

- Auto scaling provides additional performance when load grows
  - CPU utilization, network and IO bandwidth utilization thresholds**
  - Response time for new resources may not be adequate**
- As an architect you need to be aware of resource requirements for your system
  - Build and integrate the knowledge**
  - Make system self-aware and proactive**
  - Use of load balancers**



# Availability

- Failure: a regular phenomena  
**Response time for new resources may not be adequate**
- Cloud providers promise an SLA (Service level agreement)  
**99.999999...**  
**More 9s after decimal makes your system more reliable**
- As a software architect, you must assume failures  
**Build detection and correction mechanisms**
- Tactics:  
**Timeout**  
**Can not differentiate between failures and slow response**  
**Might be Long tail latency**



## Agenda for today

- 13:15 – 14:15: Cloud Architecture
- 14:30 – 14:45: Lab Assignment: I
- 14:45 – 15:45: Lab-I Assignment time
- 16:00 – 16:45: Assignment time
- 16:45 – 17:00: Wrap-up





## For Next Time



- Future computing architectures  
Edge and Quantum architectures
- Deadline Lab Assignment-I: March 8th AOE



The information in this presentation has been compiled with the utmost care,  
but no rights can be derived from its contents.